



# Software Engineer, Infrastructure

San Francisco, California, United States — Supercomputing

---

[Apply now ↗](#)

## About the Team

We believe that increasing compute is a huge lever to AI progress.

The Supercomputing team owns the entire process of building OpenAI's compute and infrastructure, which includes 1) the sourcing of hardware and system design, 2) the deployment of huge clusters using Kubernetes and Azure, 3) building the internal experiment platform for running/training the world's largest AI models.

We work at the very cutting edge of speed and scale, combining the traditions of High-Performance Computing (HPC) in a modern cloud and containerized environment.

We build some of the largest Supercomputers in the world. When our Owl cluster launched it in 2019 it would've been among the top 5 of the TOP500 supercomputers in the world. Since then we've only continued to grow. See this blog post to get a sense of what kind of challenges we solve in our day-to-day work: [Scaling Kubernetes to 7,500 Nodes](#)

You won't encounter any other organization in the world with as much compute per employee. We are a small team that moves quickly, with access to huge resources, working with a direct impact on the success of OpenAI and, by extension, the field of AI as a whole.

## About the Role

In this role, you will work with a small team of software engineers to build primarily the infrastructure and compute platform that all of the researchers at OpenAI use for their work. You do not need to be an ML/DL expert to deliver world-class infrastructure, but you do need to be able to quickly obtain a



**As a Software Engineer focused on infrastructure with the Supercomputing team, you should:**

- Know your way around bash, Terraform, Python, and/or Chef
- Have experience designing large, highly available distributed systems with Kubernetes or Mesos clusters in the range of 500+ nodes. Even better is experience with GPU workloads on those clusters.
- Have experience working with Azure or other cloud platforms such as AWS or GCP
- Have expertise debugging problems across the stack, such as networking issues, performance problems, hardware issues or memory leaks

**We think you'll find this work really exciting if you:**

- Want to help build and maintain some of the world's largest modern, cloud-based supercomputing systems
- Would enjoy working with world-class AI Research as your primary workload
- Have built large clusters but have motivation to scale beyond
- Love building large distributed, highly available systems without having to manage complex databases
- Enjoy owning things end-to-end and coming up with solutions to yet unsolved problems
- Enjoy the cycle of designing and building the next generation supercomputing cluster year after year

We estimate that someone with 3-5+ years of experience as a software engineer working on a team building and monitoring a large-scale infrastructure deployment will quickly contribute to our challenges. Any experience with high-performance computing or open-source contributions is a bonus.

We are an equal opportunity employer and do not discriminate on the basis of race, religion, national origin, gender, sexual orientation, age, veteran status, disability or any other legally protected status. Pursuant to the San Francisco Fair Chance Ordinance, we will consider qualified applicants with arrest and conviction records.

We are committed to providing reasonable accommodations to applicants with disabilities, and requests can be made via this [link](#).

[OpenAI US Applicant Privacy Policy](#)

**Compensation, Benefits and Perks**



- Mental health and wellness support
- 401(k) plan with 4% matching
- Unlimited time off and 18+ company holidays per year
- Paid parental leave (20 weeks) and family-planning support
- Annual learning & development stipend (\$1,500 per year)

### Annual Salary Range

\$200,000—\$370,000 USD

### About OpenAI

OpenAI is an AI research and deployment company dedicated to ensuring that general-purpose artificial intelligence benefits all of humanity. We push the boundaries of the capabilities of AI systems and seek to safely deploy them to the world through our products. AI is an extremely powerful tool that must be created with safety and human needs at its core, and to achieve our mission, we must encompass and value the many different perspectives, voices, and experiences that form the full spectrum of humanity.

At OpenAI, we believe artificial intelligence has the potential to help people solve immense global challenges, and we want the upside of AI to be widely shared. Join us in shaping the future of technology.

[Apply now ↗](#)



#### Research

Overview

Index

GPT-4

DALL·E 3

#### API

Overview

Data privacy

Pricing

Docs ↗



Enterprise

Try ChatGPT ↗

Blog

Careers

Charter

Security

Customer stories

Safety

---

**OpenAI © 2015–2023**

Terms & policies

Privacy policy

Brand guidelines

**Social**

Twitter

YouTube

GitHub

SoundCloud

LinkedIn

Back to top ↑



---

---

---

---